

## Data 8 Midterm Study Guide — Page 1

- Group by some *treatment* and measure some *outcome*
- Simplest setting: a *treatment group* and a *control group*
- If the *outcome* differs between these two groups, that's evidence of an *association* (or *relation*)
  - E.g., the top-tier chocolate eaters died of heart disease at a lower rate (12%) than chocolate abstainers (17%)
- If the two groups are similar in all ways but the *treatment* itself, a difference in the *outcome* is evidence of *causality*
- When a group is divided *randomly*, it's unlikely that there are systematic differences between sub-groups

Operation	Operator	Example	Value
Addition	+	2 + 3	5
Subtraction	-	2 - 3	-1
Multiplication	*	2 * 3	6
Division	/	7 / 3	2.66667
Remainder	%	7 % 3	1
Exponentiation	**	2 ** 0.5	1.41421

`more_than_1 = 2 + 3`

Name

Any expression

- Statements don't have a value; they perform an action
- An assignment statement changes the meaning of the name to the left of the = symbol
- The name is bound to a value (not an equation)

**Growth rate:** The rate of increase per unit time

- After one time unit, a quantity *x* growing at rate *g* will be  $x * (1 + g)$
- After *t* time units, a quantity *x* growing at rate *g* will be  $x * (1 + g) ** t$
- If *after* and *before* are measurements of the same quantity taken *t* time units apart, then the growth rate is  $(after/before) ** (1/t) - 1$

A **range** is an array of consecutive numbers

- `np.arange(end)`:  
An array of increasing integers from 0 up to *end*
- `np.arange(start, end)`:  
An array of increasing integers from *start* up to *end*
- `np.arange(start, end, step)`:  
A range with *step* between consecutive values

The range always includes *start* but excludes *end*

The **group** method aggregates all rows with the same value for a column into a single row in the result

- First argument: Which column to group by
- Second argument: How to combine values
  - `len` — number of grouped values (default)
  - `sum` — total of all grouped values
  - `list` — list of all grouped values

The **pivot** method aggregates by a combination of columns and organizes the result into a grid

- First argument: Which column gives the pivot columns
- Second argument: Which column gives the pivot rows
- Third argument: (Optional) Column of pivot values
- Fourth argument: (Optional) How to combine values

<code>max, min</code>	Maximum or minimum of a sequence
<code>str, float, int</code>	Convert a value to a string, number, or integer
<code>np.prod</code>	Multiply all elements together
<code>sum, np.sum</code>	Add all elements together
<code>all, np.all</code>	Test whether all elements are true values (non-zero numbers are true)
<code>any, np.any</code>	Test whether any elements are true values (non-zero numbers are true)
<code>np.count_nonzero</code>	Count the number of non-zero elements. <code>False</code> is considered zero and <code>True</code> non-zero.
<code>np.diff</code>	Difference between adjacent elements
<code>round, np.round</code>	Round each number to the nearest integer (whole number)
<code>np.cumprod</code>	A cumulative product: for each element, multiply all elements so far
<code>np.cumsum</code>	A cumulative sum: for each element, add all elements so far
<code>np.exp</code>	Exponentiate each element
<code>np.log</code>	Take the natural logarithm of each element
<code>np.abs</code>	Take the absolute value of each element
<code>np.sqrt</code>	Take the square root of each element
<code>np.sort</code>	Sort the elements

- < and > mean what you expect (less than, greater than)
- <= means "less than or equal"; likewise for >=
- == means "equal"; != means "not equal"
- Both numbers and strings can be compared, but a string cannot be compared to a number
- Comparing strings compares their alphabetical order

A table can be created using...

- `Table.read_table(url)` reads an existing table
- `Table(labels)` is a table with empty columns
- `Table()` is an empty table with no columns

A table *t* can be extended using...

- `t.with_column(label, values)`
- `t.with_columns(labels_and_values)`
- `t.with_row(row)` and `t.with_rows(rows)`

- `t.select(columns)` keeps only certain columns
- `t.drop(columns)` keeps all but certain columns
- `t.relabeled(a, b)` substitutes label *a* with *b*
- Each **column** has a label and an index (starting at 0)
  - The resulting array has one **item** per row

- `t.where(column, value)` keeps all rows containing a certain value in a column
- `t.where(condition)` keeps all rows for which the condition is true, where *condition* is a boolean array
- `t.take(row_numbers)` keeps the numbered rows
- `t.sort(column)` sorts the rows in increasing order
- Each **row** has an index (starting at 0)
  - The resulting row value has one **item** per column
  - Row items can be selected by column index or label

Tables can generate charts using these methods:

- `barh` draw a horizontal bar chart
- `hist` draw a histogram (later)
- `scatter` draw a scatter diagram (later)


## Data 8 Midterm Study Guide — Page 2

**Types of values:** All values in a column should have the same type **and** be comparable to each other in some way

- **Categorical** — Each value is from a fixed inventory
  - May or may not have an ordering
  - Categories are the same or different
  - Allows grouping by value (**group**, **pivot**, **join**)
- **Numerical** — Each value is from a fixed scale
  - Numerical measurements are ordered
  - Allows binning by value (**bin**, **hist**)

**Binning** is counting the number of numerical values that lie within ranges, called bins.

- Bins are defined by their lower bounds (inclusive)
  - The upper bound is the lower bound of the next bin
- 188, 170, 189, 163, 183, 171, 185, 168, 173, ...



A **histogram** has two defining properties:


- The bins are contiguous (though some might be empty) and are drawn to scale
- The **area** of each bar is proportional to the number of entries in the bin

A histogram drawn *on the density scale* has total area 1

Vertical axis units: Proportion / Unit on the horizontal axis

Height of a bar: (proportion of entries in bin) / (width of bin)

- A histogram of proportions of all possible outcomes of a *known* random process is called a *probability histogram*
- A histogram is a summary visualization of a *distribution*
- A histogram of proportions of actual outcomes generated by sampling is called an *empirical histogram*
- When a histogram is created from data, it is empirical



**Sampling** from a population means selecting a subset

**Probability sample:** Use a random procedure for which you *know in advance* the chance that each subset of the population will enter the sample (called the *distribution*)

**Uniform random sample** example: choose each movie at random with equal chance

Each item has equal chance of being drawn

- **With replacement:** Items can be drawn more than once, and no draw affects the chance of any other draw
- **Without replacement:** Once an item is selected, it cannot be selected again

```
def spread(values):
    return max(values) - min(values)

for i in np.arange(12):
    print(i)
```

The body is executed **for** every item in a sequence  
 The body of the statement can have multiple lines  
 The body should do something: print, assign, hist, etc.

**Law of averages:** If a chance experiment is **repeated independently** under identical conditions, in the long run:

The proportion of times that an event occurs gets closer and closer to the theoretical probability of the event

**Biased estimate:** On average across all possible samples, the estimate is either too high or too low.

- Bias creates a systematic error in one direction.
- Good estimates typically have low bias.
- The value of an estimate **varies** across samples.
- High variability makes it hard to estimate accurately.
- Good estimates typically have low variability.

**Mean (or average):** Balance point of the histogram

- **Not** the “half-way point” of the data; the mean is not the median unless the histogram is symmetric
- If the histogram is skewed, then the mean is pulled away from the median in the direction of the tail

**Standard deviation (SD)** =

root	mean	square of	deviations from	average
5	4	3	2	1

Measures roughly how far off the values are from average

“average  $\pm$  z SDs”

- z measures “how many SDs above average”
- If z is negative, the value is below average
- z is called **standard units**
- Almost all standard units are in the range (-5, 5)
- Convert a value to standard units:  $z = (\text{value} - \text{average}) / \text{SD}$

Percent in Range	All Distributions	Normal Distribution
average $\pm$ 1 SD	at least 0%	about 68%
average $\pm$ 2 SDs	at least 75%	about 95%
average $\pm$ 3 SDs	at least 88.888...%	about 99.73%

**Correlation Coefficient (r)** =

average of	product of	x in standard units	and	y in standard units
------------	------------	---------------------	-----	---------------------

Measures how clustered the scatter is around a straight line

- $-1 \leq r \leq 1$
- $r = \pm 1$  if the scatter is a perfect straight line
- r is a pure number, with no units
- r is not affected by changing units of measurement
- r is not affected by switching the horizontal and vertical axes

**Regression to the mean:** a statement about x and y pairs

- Measured in *standard units*
- Describing the deviation of x from 0 (the average of x's)
- And the deviation of y from 0 (the average of y's)

On average, y deviates from 0 less than x deviates from 0

$$\text{Regression Line } y(\text{su}) = r \times x(\text{su})$$


Correlation

In original units, the regression line has this equation:


$$\frac{\text{estimate of } y - \text{average of } y}{\text{SD of } y} = r \times \frac{\text{the given } x - \text{average of } x}{\text{SD of } x}$$

y in standard units

x in standard units



Original Units



## Data 8 Midterm Study Guide — Page 3

```
census = Table.read_table(census_url)
Census.labels      # Returns a tuple of labels
Census.num_columns # Returns an integer number of columns
Census.num_rows    # Returns an integer number of rows
census.row(5)      # Returns a row, which has items
census.column(2)   # Returns a column, which is an array
census.row(0).item(2) == census.column(2).item(0)

# Filtering on one or more conditions
census.where('AGE', 40)
census.where(census.column(2) < 40)
census.where(np.logical_or(census.column('AGE') == 18,
                           census.column('AGE') == 19))
census.where(np.logical_and(census.column('AGE') >= 18,
                           census.column('AGE') <= 19))

def age_group(age):
    if age < 2:
        return 'Baby'
    elif age < 13:
        return 'Child'
    elif age < 20:
        return 'Teen'
    else:
        return 'Adult'

population = census.with_columns([
    'Age Group', census.apply(age_group, 'AGE'),
    'Male/Female', census.column('SEX'),
    'Population', census.column('2014')])

population.group('Age Group')
# Creates a table with labels: Age Group | count

sums = population.select([0, 2]).group('Age Group', sum)
# Creates a table with labels: Age Group | Population sum

population.pivot('Age Group', 'Male/Female')
# Counts labeled: Male/Female | Adult | Baby | Child | Teen

population.join('Age Group', sums)
# Adds a column for the Population sum of each age group

top = Table.read_table('top_movies.csv')
top.take([3, 18, 100])
top.take(np.arange(0, top.num_rows, 40))
np.random.randint(3, 8) # select randomly from 3, 4, 5, 6, 7
start = np.random.randint(0, 10)
top.take(np.arange(start, top.num_rows, 10))

die = Table().with_column('Face', [1, 2, 3, 4, 5, 6])
dice_bins = np.arange(0.5, 7, 1)
die.hist(bins=dice_bins)
die.sample(10, with_replacement=True)

for i in np.arange(5):
    print(i)

# Generating a table of dice rolls
def roll_once(n):
    return np.random.randint(1, n+1)
rolls = Table(['First roll', 'Second roll'])
for i in np.arange(1000):
    row = [roll_once(6), roll_once(6)]
    rolls.append(row)
rolls.hist(bins=dice_bins)
```

<code>np.average</code> , <code>np.mean</code>	Average of values
<code>np.std</code>	Standard deviation of values
<code>np.var</code>	Variance of values (standard deviation squared)
<code>Table.append</code>	Append a row or all rows of a table.
<code>Table.append_column</code>	Appends a column to the table.
<code>Table.apply</code>	Returns an array where a function is applied to each item in a column.
<code>Table.bin</code>	Group values by bin and compute counts per bin by column.
<code>Table.column</code>	Return the values of a column as an array.
<code>Table.drop</code>	Return a table with only columns other than selected label or labels.
<code>Table.group</code>	Group rows by unique values in a column; count or aggregate other values.
<code>Table.join</code>	Generate a table with the columns of self and other, containing rows for all values of a column that appear in both tables.
<code>Table.pivot</code>	Group rows by unique values in two columns; count or aggregate values from a third column.
<code>Table.relabel</code>	Change the label of a column
<code>Table.relabeled</code>	Return a new table with a label changed.
<code>Table.sample</code>	Returns a new table where k rows are randomly sampled from the original table.
<code>Table.select</code>	Return a table with only the selected columns.
<code>Table.sort</code>	Return a table of rows sorted according to the values in a column.
<code>Table.take</code>	Return a table of the rows taken by index.
<code>Table.where</code>	Return a table of the rows for which the column matches a value or is a non-zero value.
<code>Table.with_column</code> <code>Table.with_columns</code>	Return a table with an additional or replaced column or columns.
<code>Table.with_row</code> <code>Table.with_rows</code>	Return a table with an additional row or rows.

```
annual_growth_rate = 0.035
```

```
ten_year_growth_rate = (1 + annual_growth_rate) ** 10 - 1
```

```
ten_year_growth_rate # evaluates to: 0.410598760621121
```

```
baseline_high = 14.48
```

```
highs = np.array([baseline_high - 0.880, baseline_high - 0.093,
                  baseline_high + 0.105, baseline_high + 0.684])
```

```
highs # evaluates to: array([ 13.6 , 14.387, 14.585, 15.164])
(9/5) * highs + 32 # array([ 56.48 , 57.89, 58.253, 59.295])
```

```
def percent(x):
```

```
    if x < 1e-8:
```

```
        return 0.0
```

```
    elif x < 0.00005:
```

```
        return 100 * x
```

```
    else:
```

```
        return round(100 * x, 2)
```

```
percent(1/6) # returns 16.67
```

```
percent(1/6000) # returns 0.02
```

```
percent(1/60000) # returns 0.0016666666666666668
```

```
percent(1/60000000000) # returns 0.0
```

```
Table(['Odd', 'Even']).with_row([3, 4])
```

```
Table(['Odd', 'Even']).with_rows([3, 4], [5, 6], [7, 8])
```

```
Table().with_column('Odd', [3, 5, 7])
```

```
Table().with_columns(['Odd', [3, 5, 7], 'Even', [4, 6, 8]])
```

```

def standard_units(any_numbers):
    return (any_numbers - np.mean(any_numbers))/np.std(any_numbers)

def correlation(t, x, y):
    return np.mean(standard_units(t.column(x))*standard_units(t.column(y)))

def slope(table, x, y):
    r = correlation(table, x, y)
    return r * np.std(table.column(y))/np.std(table.column(x))

def intercept(table, x, y):
    a = slope(table, x, y)
    return np.mean(table.column(y)) - a * np.mean(table.column(x))

def fit(table, x, y):
    a = slope(table, x, y)
    b = intercept(table, x, y)
    return a * table.column(x) + b

# Regression effect
heights = Table.read_table('heights.csv')
six_foot_fathers = heights.where(np.round(heights.column('father')) == 72)
six_foot_fathers.hist('son', bins=np.arange(55.5, 80, 1))
np.count_nonzero(six_foot_fathers.column('son') < 72) / six_foot_fathers.num_rows
np.mean(six_foot_fathers.column('son'))

# Graph of averages
fathers = heights.column('father')
sons = heights.column('son')
heights_su = Table().with_columns([
    'father (su)', standard_units(fathers),
    'son (su)', standard_units(sons)])
rounded = heights_su.with_column('father (su)', np.round(father_su))
rounded.join('father (su)', rounded.group(0, np.average)).scatter(0, s=80)

# Regression model
def signal_and_noise(x, r):
    return r * x + np.random.normal() * (1-r**2)**0.5

def regression_model(r, sample_size):
    pairs = Table(['x', 'y'])
    for i in np.arange(sample_size):
        x = np.random.normal()
        y = signal_and_noise(x, r)
        pairs.append([x, y])
    return pairs

regression_model(1/2, 1000).scatter('x', 'y')

def compare(true_r, sample_size):
    pairs = regression_model(true_r, sample_size)
    estimated_r = correlation(pairs, 'x', 'y')
    pairs.scatter('x', 'y', fit_line=True, s=10)
    plt.plot([-3, 3], [-3 * true_r, 3 * true_r], color='g', lw=4)
    print("The true r is ", true_r, " and the estimated r is ", estimated_r)

compare(0.5, 1000) # Compare true line to regression line

```



## Data 8 Final Study Guide — Page 1

Higher-order function: A function that returns a function

```
def <name>(x):  
    def <inner name>(y):  
        return <some expression involving x and y>  
    return <inner name>
```

The returned *function* takes one argument (*y*) but computes its result from two values (*x* and *y*).

The regression line is the one that minimizes the (root) mean squared error of a collection of paired values

- The slope and intercept are unique for linear regression

```
def mean_squared_error(table, x, y):  
    def for_line(a, b):  
        estimate = (a * table.column(x) + b)  
        return np.average((table.column(y) - estimate) ** 2)  
    return for_line
```

```
long = faithful.where('eruptions', are.above(3))
```

```
mse_long = mean_squared_error(long, 1, 0)
```

```
a, b = minimize(mse_long) # regression line is a * x + b
```

Simple regression: one input → one output

Multiple regression: many inputs → one output

$$\text{GPA} = a_{\text{days}} * \text{days} + a_{\text{contributions}} * \text{contributions} + b$$

We could find *a*'s and *b* by minimizing mean squared error

**Inference:** Making conclusions from random samples

**Population:** The entire set that is the subject of interest

**Parameter:** A quantity computed for the entire population

**Sample:** A subset of the population

In a **Random Sample**, we know the chance that any subset of the population will enter the sample, in advance

**Statistic:** A quantity computed for a particular sample

Estimation is a process with a random outcome

Population (fixed) → Sample (random) → Statistic (random)

A 95% **Confidence Interval** is an interval constructed so that it will contain the parameter for 95% of samples

For a particular sample, the interval either contains the parameter or it doesn't; the process works 95% of the time

**Resampling:** When we wish we could sample again from the population, instead sample from the sample

The 80th percentile is the value in a set that is at least as large as 80% of the elements in the set

For *s* = [1, 7, 3, 9, 5], `percentile(80, s)` is 7


The 80th percentile is ordered element 4:  $(80/100) * 5$

For a percentile that does not exactly correspond to an element, take the next greater element instead

Percentile      Size of set

`percentile(10, s)` is 1      `percentile(20, s)` is 1  
`percentile(21, s)` is 3      `percentile(40, s)` is 3

If an interval **around the parameter** contains the estimate, then a (reflected) interval of the same width **around the estimate** contains the parameter (and vis versa)



**Inferential idea:** The variability of a statistic for different samples is similar to the variability of that statistic for resamples. Computing a confidence interval for an estimate from a sample:

- Collect a random **sample**
- Compute your **estimate** (e.g., **sample** average)
- **Resample K samples** from the **sample**, with replacement.
  - Compute the same statistic for each **resampled sample**
  - Take **percentiles** of the **deviations** around the **estimate**
  - 95%: (**estimate** - 97½% deviation, **estimate** - 2½% deviation)

## K-nearest Neighbor classifier

**Goal:** Using a set of labeled examples, define a classifier that can predict the label of an unlabeled example

**Example:** A set of features (i.e., a row in a table of features)

**Data preparation:** Randomly divide all available examples into three sets: training, validation, and test

**Classification:** Find the K examples in the training set that are most similar to the example and average the labels

**Accuracy:** The proportion of examples in a set (e.g., the test set or validation set) that are classified with the correct label

```
def distance(pt1, pt2):  
    tot = 0  
    for i in range(len(pt1)):  
        tot = tot + (pt1[i] - pt2[i])**2  
    return math.sqrt(tot)  
  
def computetablewithdists(training, p):  
    dists = np.zeros(training.num_rows)  
    attributes = training.drop('Class')  
    for i in np.arange(training.num_rows):  
        dists[i] = distance(attributes.row(i), p)  
    return training.with_column('Distance', dists)  
  
def closest(training, p, k):  
    withdists = computetablewithdists(training, p)  
    sortedbydist = withdists.sort('Distance')  
    topk = sortedbydist.take(np.arange(k))  
    return topk  
  
def majority(topk):  
    if topk.where('Class', 1).num_rows > topk.where('Class', 0).num_rows:  
        return 1  
    else:  
        return 0  
  
def classify(training, p, k):  
    closestk = closest(training, p, k)  
    topkclasses = closestk.select('Class')  
    return majority(topkclasses)  
  
def evaluate_accuracy(training, test, k):  
    testattrs = test.drop('Class')  
    numcorrect = 0  
    for i in range(test.num_rows):  
        # Run the classifier on the ith patient in the test set  
        c = classify(training, testattrs.rows[i], k)  
        # Was the classifier's prediction correct?  
        if c == test.column('Class').item(i):  
            numcorrect = numcorrect + 1  
    return numcorrect / test.num_rows
```

## Feature Selection Using the Validation Set

**Goal:** Select features and K to maximize accuracy

**Overfitting:** Accuracy is higher on the set used for feature selection than another set because of peculiarities of the set

**Selecting Features:** Maintain a set of selected features, which starts out empty; Each added feature is the feature that improves validation set accuracy the most, in combination with the features already added

```
def confidence_interval_95_percent(sample, label, f):  
    deviations = Table(['Resample #', 'Deviation'])  
    n = sample.num_rows  
    stat = f(sample.column(label))  
    for i in np.arange(400):  
        resample = sample.sample(n, with_replacement=True)  
        dev = f(resample.column(label)) - stat  
        deviations.append([i, dev])  
    return (stat - percentile(97.5, deviations.column(1)),  
            stat - percentile(2.5, deviations.column(1)))
```

## Data 8 Final Study Guide — Page 2

### Total Variation Distance (TVD):

- For each category, compute the difference in proportions between two distributions
- Take the absolute value of each difference
- Sum and divide by 2

```
def total_variation_distance(column, other):  
    return sum(np.abs(column - other)) / 2
```

`np.random` is a collection of sampling procedures

```
np.random.randint(lower, upper)
```

An integer lower bound (inclusive)

An integer upper bound (exclusive)

```
np.random.multinomial(n, distribution)
```

The number of samples to draw

An array or list containing the chance for each category

### Does a sample look like a random sample from a distribution?

- E.g., Do the counts of races on a jury panel look like a random sample from the proportion of races in a county?
- Sample many times from the distribution and compute the TVD of each sample from the original distribution
  - The result is called an *empirical distribution* of a statistic
  - It approximates the *sampling distribution* of the statistic
- Compare the observed TVD between the sample and the original distribution to this *empirical distribution*

```
def empirical_distribution(table, label, size, k, f):  
    stats = Table(['Sample #', 'Statistic'])  
    for i in np.arange(k):  
        sample = table.sample_from_distribution(label, size)  
        statistic = f(sample)  
        stats.append([i, statistic])  
    return stats
```

### Step 1: The Hypotheses

- A test chooses between two views of how data were generated
- Null hypothesis* proposes that data were generated at random
- Alternative hypothesis* proposes some effect other than chance

### Step 2: The Test Statistic

- A value that can be computed for the data and for samples

### Step 3: The Sampling Distribution of the Test Statistic

- What the test statistic might be if the null hypothesis were true
- Approximate the sampling distribution by an empirical distribution

### Step 4: Resolve choice between null and alternative hypotheses

- Compare observed test statistic to its empirical distribution under the null hypothesis
- If the observed value is **consistent** with the distribution, then the test *does not* support the alternative hypothesis

**P-Value:** The chance, under the null hypothesis, that the test statistic is equal to the value that was observed or is even further in the direction of the alternative.

**Statistically Significant:** The P-value is less than 5%

**Highly Statistically Significant:** The P-value is less than 1%

### Comparing samples: Do two groups differ in some attribute?

- E.g.: Among U.S. men who are in heterosexual relationships, is there an association between marriage and employment status?
- Null hypothesis:** In the United States, the distribution of employment status among married men is the same as among unmarried men who live with their partners.
- Inferential Idea:** If marital status and employment status were not connected in any way, then we could replicate the sampling process by replacing each man's employment status by a randomly picked employment status from among all the men
  - Permute (shuffle) the outcome column K times
  - A shuffled table pairs each example with a random outcome
  - Compute the total variation distance between the two treatment conditions to create an empirical distribution

### Bootstrap A/B test using means

- $n$  individuals in Category A;  $m$  in Category B
- Under the null hypothesis, samples in both categories are drawn from the same underlying distribution.
- So draw  $(n+m)$  times at random with replacement from the entire sample. Assign  $n$  of the draws to Category A and  $m$  to Category B.
- Find the difference between the means of these two new groups.
- Repeat the last two steps many times and compare with the difference between the two original sample means.

We observed a positive slope and used it to make our predictions.



But what if the scatter plot got its positive slope just by chance?



What if the true line is actually FLAT?

### Bootstrap the scatter plot & find the slope of the regression line through the bootstrapped plot many times.

- Draw the empirical histogram of all the resampled slopes.
- Get the "middle 95%" interval: that's an approximate 95% confidence interval for the slope of the true line.
- Null hypothesis:** The slope of the true line is 0.
  - Construct a bootstrap confidence interval for the true slope.
  - If the interval doesn't contain 0, reject the null hypothesis.
  - If the interval does contain 0, there isn't enough evidence to reject the null hypothesis.


- Residual:** Difference between  $y$  and regression line height at  $x$ .
- Regression Model:**  $y$  is a linear function of  $x$  + normal "noise"
- Residual plot looks like a formless "noise" cloud under this model
- Average of residuals is always 0 for any scatter diagram

$$|r| = \frac{\text{SD of fitted values of } y}{\text{SD of observed values of } y} \quad \text{SD of residuals} = \sqrt{1 - r^2} \times \text{SD of } y$$

- Experiment:** An occurrence with an uncertain outcome
- Outcome:** The result of an experiment
- Sample Space:** The set of all possible outcomes for the experiment
- Event:** A subset of possible outcomes that have some property of interest
- Probability:** The proportion of experiments for which the event occurs
- Distribution:** The probability of all events

**Multinomial distributions** have a finite sample space

- Distribution with a fixed number of possible outcomes
- All outcomes are mutually exclusive for an experiment
- The sum of the chances of the outcomes is 1
- The probability of an event is the sum of the chances of the outcomes in which the event occurs



**Conditional Probability:** The chance of an event B **given** an event A  
Both events describe the same experiment & sample space

- Find all outcomes in which A occurs
- Divide the chances of A's outcomes by  $P(A)$
- Find all outcomes in which B also occurs
- Sum the updated chances of these outcomes:  $P(B | A) = P(A, B) / P(A)$
- Bayes' Rule:  $P(A | B) = P(B | A) * P(A) / P(B)$

Conditional Distribution

repetitions = 500

```
tvs = Table().with_column("TVD between married and partnered men", [])
```

```
for i in np.arange(repetitions):
```

```
# Construct a permuted table
```

```
shuffled = males.select('Employment Status').sample()
```

```
combined = Table().with_columns([
```

```
    "Marital Status", males.column('Marital Status'),
```

```
    "Employment Status", shuffled.column('Employment Status')])
```

```
employment_shuffled = combined.pivot('Marital Status', 'Employment Status')
```

```
# Compute TVD
```

```
e_s = employment_shuffled
```

```
married = e_s.column('married') / sum(e_s.column('married'))
```

```
partner = e_s.column('partner') / sum(e_s.column('partner'))
```

```
permutation_tvd = 0.5 * sum(abs(married - partner))
```

```
tvs.append([permutation_tvd])
```

x	y	
a	x	Randomly permute (shuffle) column y
a	x	
b	y	
b	z	
b	z	
b	y	

## Exam Scratch Paper

## Exam Scratch Paper



## Exam Scratch Paper

## Exam Scratch Paper